

Using Cross Validation

In this exercise, you will use cross-validation to optimize parameters for a regression model.

Prepare the Data

First, import the libraries you will need and prepare the training and test data:

In []:

```
In [1]: # Import Spark SQL and Spark ML libraries
from pyspark.sql.types import *
from pyspark.sql.functions import *

from pyspark.ml import Pipeline
from pyspark.ml.regression import LinearRegression
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
from pyspark.ml.evaluation import RegressionEvaluator

# Load the source data
csv = spark.read.csv('wasb://data/flights.csv', inferSchema=True, header=True)

# Select features and Label
data = csv.select("DayofMonth", "DayOfWeek", "OriginAirportID", "DestAirportID",

# Split the data
splits = data.randomSplit([0.7, 0.3])
train = splits[0]
test = splits[1].withColumnRenamed("label", "trueLabel")
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI
7	application_1613740567063_0010	pyspark	idle	Link (http://hn1-rayspa.pk5nseqskssvekfuqghr2123nc.ix.int)

SparkSession available as 'spark'.

Define the Pipeline

Now define a pipeline that creates a feature vector and trains a regression model

```
In [2]: # Define the pipeline
assembler = VectorAssembler(inputCols = ["DayofMonth", "DayOfWeek", "OriginAirportID"])
lr = LinearRegression(labelCol="label",featuresCol="features")
pipeline = Pipeline(stages=[assembler, lr])
```

Tune Parameters

You can tune parameters to find the best model for your data. To do this you can use the **CrossValidator** class to evaluate each combination of parameters defined in a **ParameterGrid** against multiple *folds* of the data split into training and validation datasets, in order to find the best performing parameters. Note that this can take a long time to run because every parameter combination is tried multiple times.

```
In [3]: paramGrid = ParamGridBuilder().addGrid(lr.regParam, [0.3, 0.01]).addGrid(lr.maxIter, [10, 20])
cv = CrossValidator(estimator=pipeline, evaluator=RegressionEvaluator(), estimatorParamMaps=paramGrid.build(),
model = cv.fit(train)
```

Test the Model

Now you're ready to apply the model to the test data.

```
In [4]: prediction = model.transform(test)
predicted = prediction.select("features", "prediction", "trueLabel")
predicted.show()
```

features	prediction	trueLabel
[1.0,1.0,10140.0,...]	73.62671581440335	94
[1.0,1.0,10140.0,...]	-8.892516123547134	-14
[1.0,1.0,10140.0,...]	-6.882329878216394	-11
[1.0,1.0,10140.0,...]	-5.877236755551024	-12
[1.0,1.0,10140.0,...]	-4.872143632885654	-11
[1.0,1.0,10140.0,...]	17.239905065752485	19
[1.0,1.0,10140.0,...]	31.31120878306766	41
[1.0,1.0,10140.0,...]	-5.884891239614667	-9
[1.0,1.0,10140.0,...]	-5.884891239614667	-5
[1.0,1.0,10140.0,...]	-1.8645187489531878	-1
[1.0,1.0,10140.0,...]	-0.8594256262878179	2
[1.0,1.0,10140.0,...]	-13.927027945312835	-13
[1.0,1.0,10140.0,...]	-5.886282963989875	-2
[1.0,1.0,10140.0,...]	-4.881189841324505	-9
[1.0,1.0,10140.0,...]	-2.8710035959937654	-3
[1.0,1.0,10140.0,...]	-15.156652600511714	-28
[1.0,1.0,10140.0,...]	-12.141373232515605	-17
[1.0,1.0,10140.0,...]	838.1674085423874	812
[1.0,1.0,10140.0,...]	-8.265508122813193	3
[1.0,1.0,10140.0,...]	-6.255321877482453	-17

only showing top 20 rows

Examine the Predicted and Actual Values

You can plot the predicted values against the actual values to see how accurately the model has predicted. In a perfect model, the resulting scatter plot should form a perfect diagonal line with each predicted value being identical to the actual value - in practice, some variance is to be expected. Run the cells below to create a temporary table from the **predicted** DataFrame and then retrieve the predicted and actual label values using SQL. You can then display the results as a scatter plot, specifying `-` as the function to show the unaggregated values.

```
In [5]: predicted.createOrReplaceTempView("regressionPredictions")
```

```
In [6]: %%sql  
SELECT trueLabel, prediction FROM regressionPredictions
```

Type: Table Pie Scatter Line Area Bar

Encoding:

X	trueLabel
Y	prediction

Func. -

Log scale X

Log scale Y

Retrieve the Root Mean Square Error (RMSE)

There are a number of metrics used to measure the variance between predicted and actual values. Of these, the root mean square error (RMSE) is a commonly used value that is measured in the same units as the predicted and actual values - so in this case, the RMSE indicates the average number of minutes between predicted and actual flight delay values. You can use the **RegressionEvaluator** class to retrieve the RMSE.

```
In [7]: evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction",  
rmse = evaluator.evaluate(prediction)  
print "Root Mean Square Error (RMSE):", rmse
```

Root Mean Square Error (RMSE): 13.246528636

```
In [ ]:
```